

SAMPLE
CHAPTER

covers Spring 4

Spring IN ACTION

FOURTH EDITION

Craig Walls





Spring in Action
by Craig Walls

Chapter 12

brief contents

PART 1 CORE SPRING1

- 1 ■ Springing into action 3
- 2 ■ Wiring beans 32
- 3 ■ Advanced wiring 64
- 4 ■ Aspect-oriented Spring 97

PART 2 SPRING ON THE WEB..... 129

- 5 ■ Building Spring web applications 131
- 6 ■ Rendering web views 164
- 7 ■ Advanced Spring MVC 194
- 8 ■ Working with Spring Web Flow 219
- 9 ■ Securing web applications 244

PART 3 SPRING IN THE BACK END 279

- 10 ■ Hitting the database with Spring and JDBC 281
- 11 ■ Persisting data with object-relational mapping 305
- 12 ■ Working with NoSQL databases 327
- 13 ■ Caching data 362
- 14 ■ Securing methods 379

PART 4 INTEGRATING SPRING.....391

- 15 ■ Working with remote services 393
- 16 ■ Creating REST APIs with Spring MVC 416
- 17 ■ Messaging in Spring 452
- 18 ■ Messaging with WebSocket and STOMP 485
- 19 ■ Sending email with Spring 511
- 20 ■ Managing Spring beans with JMX 523
- 21 ■ Simplifying Spring development with Spring Boot 540

12

Working with NoSQL databases

This chapter covers

- Writing repositories backed by MongoDB and Neo4j
- Persisting data across multiple data stores
- Working with Spring and Redis

In his autobiography, Henry Ford is famous for having written “Any customer can have a car painted any color that he wants so long as it is black.”¹ Some say this statement was arrogant and bull-headed. Others may think he was showing a bit of humor. The reality, however, may be found in the fact that at the time his biography was published, he was cutting costs by using a quick-drying paint that was only available in black.

Paraphrasing Ford’s famous quote and applying it to database choice, we’ve been told for years that we can have any database we want, as long as it’s a relational database. Relational databases have had a near-monopolistic hold on application development for a very long time.

¹ Henry Ford and Samuel Crowther, *My Life and Work* (Garden City, New York: Garden City Publishing Company, 1922).

But that hold is weakening now that some serious contenders have entered the database space. The so-called “NoSQL” databases are making inroads into production applications everywhere as we recognize that there’s no one-size-fits-all database. We now have a greater choice and can choose the best database for the problem we’re trying to solve.

Over the past couple of chapters, we’ve focused on relational databases, starting with Spring’s JDBC support and then object-relational mapping. In the previous chapter, specifically, we looked at Spring Data JPA, one of several projects under the Spring Data umbrella project. We saw how Spring Data JPA can make working with JPA more pleasant by automatically generating repository implementations at runtime.

Spring Data also supports several NoSQL databases, including MongoDB, Neo4j, and Redis. This not only includes support for automatic repositories, but also template-based data access and mapping annotations. In this chapter, we’re going to see how to write repositories that work with non-relational, NoSQL databases. We’ll start with Spring Data MongoDB to see how to write repositories that deal with document-based data.

12.1 Persisting documents with MongoDB

Some kinds of data are best represented as *documents*. That is, rather than spread the data across multiple tables, nodes, or entities, it may make more sense to collect the information into denormalized structures (known as documents). Although two or more of these documents may be related to each other, generally documents are standalone entities. Databases that are finely tuned to work with documents in this way are known as document databases.

For example, suppose that you’re writing an application that captures a college student’s transcript. You’ll need to be able to retrieve transcripts given a student’s name or perhaps search across the transcripts for some common properties. But each student is evaluated individually, so it isn’t necessary for any two transcripts to be related to each other. Although a relational database schema could be (and probably has been) designed to capture this transcript data, perhaps a document database is a better choice.

What document databases aren’t good for

Knowing when to use a document database is important. But it’s also important to know when document databases don’t make sense. Document databases aren’t general-purpose databases and they have a very narrow set of problems that they address well.

Document databases aren’t well-tuned for storing data where there’s any significant degree of relationships. A social network, for example, represents how different users of an application relate to each other and isn’t best kept in a document database. Even though it’s not impossible to store relation-rich data in a document database, you’ll find more challenge than benefit in doing so.

The domain of the Spitttr application isn't a good fit for a document database. In this chapter, we'll look at MongoDB in the context of a purchase order system.

MongoDB is one of the most popular open source document databases available. Spring Data MongoDB brings MongoDB to Spring applications in three ways:

- Annotations for object-to-document mapping
- Template-based database access with `MongoTemplate`
- Automatic runtime repository generation

We've already looked at how Spring Data JPA enabled automatic repository generation for JPA-based data access. Spring Data MongoDB offers the same feature for MongoDB-based data access.

Unlike Spring Data JPA, however, Spring Data MongoDB also offers annotations to map Java objects to documents. (Spring Data JPA doesn't need to offer such annotations for JPA because the JPA specification itself defines object-to-relational mapping annotations.) Moreover, Spring Data MongoDB provides for template-based MongoDB data access for several common document manipulation tasks.

Before we can use any of these features, however, we'll need to configure Spring Data MongoDB.

12.1.1 Enabling MongoDB

In order to effectively work with Spring Data MongoDB, you're going to need a few essential beans in your Spring configuration. First, you'll need to configure a `MongoClient` bean to be able to access the MongoDB database. You'll also need a `MongoTemplate` bean to be able to perform template-based data access against the database. Optionally, but desirably, you'll want to enable Spring Data MongoDB's automatic repository generation.

The following listing shows how to write a simple Spring Data MongoDB configuration class that addresses these needs.

Listing 12.1 An essential configuration for Spring Data MongoDB

```
package orders.config;
import org.springframework.context.annotation.Bean;
import org.springframework.context.annotation.Configuration;
import org.springframework.data.mongodb.core.MongoFactoryBean;
import org.springframework.data.mongodb.core.MongoOperations;
import org.springframework.data.mongodb.core.MongoTemplate;
import org.springframework.data.mongodb.repository.config.
    EnableMongoRepositories;

import com.mongodb.Mongo;

@Configuration
@EnableMongoRepositories(basePackages="orders.db")
public class MongoConfig {
```

← Enable MongoDB
repositories

```

@Bean
public MongoFactoryBean mongo() {
    MongoFactoryBean mongo = new MongoFactoryBean();
    mongo.setHost("localhost");
    return mongo;
}

@Bean
public MongoOperations mongoTemplate(Mongo mongo) {
    return new MongoTemplate(mongo, "OrdersDB");
}
}

```

← **MongoClient bean**

← **MongoTemplate bean**

As you'll recall from the previous chapter, you enabled Spring Data's automatic JPA repository generation with the `@EnableJpaRepositories` annotation. Similarly, the `@EnableMongoRepositories` annotation does the same thing for MongoDB.

In addition to `@EnableMongoRepositories`, listing 12.1 also includes two `@Bean` methods. The first `@Bean` method uses `MongoFactoryBean` to declare a `Mongo` instance. This bean will bridge Spring Data MongoDB to the database itself (not unlike what a `DataSource` does when working with a relational database). Although you could create an instance of `Mongo` directly with `MongoClient`, you'd be forced to deal with the `UnknownHostException` that's thrown from `MongoClient`'s constructor. It's easier to use Spring Data MongoDB's `MongoFactoryBean` here. As a factory bean, `MongoFactoryBean` will construct an instance of `Mongo` for you, without you needing to worry much about `UnknownHostException`.

The other `@Bean` method declares a `MongoTemplate` bean. It's constructed giving it a reference to the `Mongo` instance created by the other bean method and the name of the database. In a moment, you'll see how to use `MongoTemplate` to query the database. Even if you never use `MongoTemplate` directly, you'll need this bean because the automatically generated repositories will use it under the covers.

Rather than declare those beans directly, the configuration class could extend `AbstractMongoConfiguration` and override its `getDatabaseName()` and `mongo()` methods. The following listing shows how.

Listing 12.2 Enabling Spring Data MongoDB with the `@EnableMongoRepositories`

```

package orders.config;
import org.springframework.context.annotation.Configuration;
import org.springframework.data.mongodb.config.
    AbstractMongoConfiguration;
import org.springframework.data.mongodb.repository.config.
    EnableMongoRepositories;

import com.mongodb.Mongo;
import com.mongodb.MongoClient;

@Configuration
@EnableMongoRepositories("orders.db")
public class MongoConfig extends AbstractMongoConfiguration {

    @Override

```



```
protected String getDatabaseName() {           ← Specify database name
    return "OrdersDB";
}

@Override
public Mongo mongo() throws Exception {        ← Create a Mongo client
    return new MongoClient();
}

}
```

This new configuration class is equivalent to the one in listing 12.1, albeit marginally simpler. The most noticeable difference is that this configuration doesn't directly declare a `MongoTemplate` bean, although one is implicitly created. Instead, you override `getDatabaseName()` to provide the name of the database. The `mongo()` method still creates an instance of `MongoClient`, but because it throws `Exception`, you can work with `MongoClient` directly without working with `MongoFactoryBean`.

As it stands, either listing 12.1 or 12.2 provide a working configuration for Spring Data MongoDB. That is, as long as the MongoDB server is running on localhost. If your MongoDB server is running on a different server, you can specify that when you create `MongoClient`:

```
public Mongo mongo() throws Exception {
    return new MongoClient("mongodbserver");
}
```

It's also possible that your MongoDB server is listening on a port other than the default (27017). In that case, you should also specify the port when creating `MongoClient`:

```
public Mongo mongo() throws Exception {
    return new MongoClient("mongodbserver", 37017);
}
```

And if your MongoDB server is running in a production setting, I'd hope that you have authentication enabled. In that case, you'll need to provide your application's credentials in order to access the database. Accessing an authenticated MongoDB server is a bit more involved, as you can see in the next listing.

Listing 12.3 Creating a MongoClient to access an authenticated MongoDB server

```
@Autowired
private Environment env;

@Override
public Mongo mongo() throws Exception {
    MongoCredential credential =
        MongoCredential.createMongoCRCredential(
            env.getProperty("mongo.username"),
            "OrdersDB",
            env.getProperty("mongo.password").toCharArray());

    return new MongoClient(
        new ServerAddress("localhost", 37017),
        Arrays.asList(credential));
}
```

In order to access an authenticated MongoDB server, `MongoClient` must be instantiated with a list of `MongoCredentials`. In listing 12.3 a single `MongoCredential` is created for that purpose. In order to keep the credential details out of the configuration class, they're resolved from the injected `Environment`.

For what it's worth, Spring Data MongoDB can also be configured in XML. As you should know by now, I favor the Java configuration option. But if you've got a fondness for XML configuration, the following listing gives an example of how to configure Spring Data MongoDB using the `mongo` configuration namespace.

Listing 12.4 Spring Data MongoDB offers an XML configuration option

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:mongo="http://www.springframework.org/schema/data/mongo"
  xsi:schemaLocation="
    http://www.springframework.org/schema/data/mongo
    http://www.springframework.org/schema/data/mongo/spring-mongo.xsd
    http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd">

  <mongo:repositories base-package="orders.db" />
  <mongo:mongo />

  <bean id="mongoTemplate"
    class="org.springframework.data.mongodb.core.MongoTemplate">
    <constructor-arg ref="mongo" />
    <constructor-arg value="OrdersDB" />
  </bean>

</beans>
```

Declare MongoClient →

Declare mongo namespace ←

Enable repository generation ←

Create MongoTemplate bean ←

Now that Spring Data MongoDB has been configured, you're almost ready to start using it to save and retrieve documents. But first, you'll need to map your Java domain types for document persistence using Spring Data MongoDB's object-to-document mapping annotations.

12.1.2 Annotating model types for MongoDB persistence

When working with JPA, you had to map your Java entity types to relational tables and columns. The JPA specification provides for several annotations to support object-to-relational mapping, and some JPA implementations, such as Hibernate, add their own mapping annotations as well.

MongoDB, however, doesn't come with its own object-to-document mapping annotations. Spring Data MongoDB seized the opportunity to fill that gap with a handful of annotations that you can use to map your Java types to MongoDB documents. Table 12.1 describes these annotations.

Table 12.1 Spring Data MongoDB annotations for object-to-document mapping

Annotation	Description
@Document	Identifies a domain object to be mapped to a MongoDB document
@Id	Indicates that a field is the ID field
@DbRef	Indicates that a field is intended to reference another document, possibly in another database
@Field	Defines custom metadata for a document field
@Version	Identifies a property to be used as a version field

The @Document and @Id annotations are analogous to JPA's @Entity and @Id annotations. You'll use these two annotations often and on every Java type that will be stored as a document in the MongoDB database. For example, the next listing shows how you might annotate an Order class to be persisted in MongoDB.

Listing 12.5 Spring Data MongoDB annotations map Java types to documents.

```
package orders;
import java.util.Collection;
import java.util.LinkedHashSet;
import org.springframework.data.annotation.Id;
import org.springframework.data.mongodb.core.mapping.Document;
import org.springframework.data.mongodb.core.mapping.Field;

@Document                                     ← This is a document
public class Order {

    @Id
    private String id;                        ← Designate the ID

    @Field("client")
    private String customer;                  ← Override the default field name

    private String type;

    private Collection<Item> items = new LinkedHashSet<Item>();

    public String getCustomer() {
        return customer;
    }

    public void setCustomer(String customer) {
        this.customer = customer;
    }

    public String getType() {
        return type;
    }

    public void setType(String type) {
        this.type = type;
    }
}
```

```

public Collection<Item> getItems() {
    return items;
}

public void setItems(Collection<Item> items) {
    this.items = items;
}

public String getId() {
    return id;
}
}

```

As you can see, `Order` is annotated with `@Document`, enabling it to be persisted using `MongoTemplate`, an automatically generated repository, or both. Its `id` property is annotated with `@Id` to designate it as the ID of the document. In addition, the `customer` property is annotated with `@Field` so that when the document is persisted, the `customer` property will be mapped to a field named `client`.

Notice that no other properties are annotated. Unless they're marked as transient, all fields of the Java object will be persisted as fields of the document. And unless otherwise indicated with `@Field`, the document fields will have the same names as their Java property counterparts.

Also, take note of the `items` property. It's clearly a collection of line items in this order. In a traditional relational database setting, those items would probably be kept in a separate database table, referenced with a foreign key, and the `items` field might be annotated for JPA with `@OneToMany`. But here that's not the case.

As I said earlier, documents can be related to other documents, but that's not what document databases are especially good at. In the case of the relationship between a purchase order and its line items, the line items are merely a nested part of the same order document (as shown in figure 12.1). Therefore, there's no need for any annotations to designate the relationship. In fact, the `Item` class itself isn't annotated at all:

```

package orders;

public class Item {

    private Long id;
    private Order order;
    private String product;
    private double price;
    private int quantity;
}

```




Figure 12.1 Documents represent related but denormalized data. Related concepts (such as items of an order) are embedded in the top-level document.

```

public Order getOrder() {
    return order;
}

public String getProduct() {
    return product;
}

public void setProduct(String product) {
    this.product = product;
}

public double getPrice() {
    return price;
}

public void setPrice(double price) {
    this.price = price;
}

public int getQuantity() {
    return quantity;
}

public void setQuantity(int quantity) {
    this.quantity = quantity;
}

public Long getId() {
    return id;
}
}

```

It's not necessary to annotate `Item` with `@Document`, nor is it necessary to annotate one of its fields with `@Id`. That's because you'll never persist an `Item` as an independent document. It will always be a member of the `Order` document's `Item` list and a nested element in that document.

Of course, you could annotate one of `Item`'s properties with `@Field` if you wanted to dictate how that field should be stored in the document. It just wasn't necessary to do so in this example.

Now we have a Java domain type annotated for MongoDB persistence. Let's see how you can use `MongoTemplate` to store a few of them.

12.1.3 Accessing MongoDB with MongoTemplate

You've already configured a `MongoTemplate` bean, either explicitly or by extending `AbstractMongoConfiguration` in your configuration class. All you need to do is inject it wherever it will be used:

```

@Autowired
MongoOperations mongo;

```

Notice that here I'm injecting `MongoTemplate` into a property whose type is `MongoOperations`. `MongoOperations` is an interface that `MongoTemplate` implements, and

it's good form to not work with the concrete implementation directly, especially when it's injected.

`MongoOperations` exposes several useful methods for working with a MongoDB document database. It'd be impossible for us to discuss all of them here, but we can take a look at a few of the most commonly used operations, such as counting how many documents are in a document collection. Using the injected `MongoOperations`, you get the order collection and then call `count()` to get a count:

```
long orderCount = mongo.getCollection("order").count();
```

Now let's suppose you want to save a new order. To do that, you can call the `save()` method:

```
Order order = new Order();  
... // set properties and add line items  
mongo.save(order, "order");
```

The first parameter to `save()` is the newly created `Order`; the second is the name of the document store to save it to.

You can also look up an order by its ID by calling `findById()`:

```
String orderId = ...;  
Order order = mongo.findById(orderId, Order.class);
```

More advanced queries require that you construct a `Query` object and pass it to the `find()` method. For example, to find all orders whose `client` field is equal to "Chuck Wagon", you can use this code:

```
List<Order> chucksOrders = mongo.find(Query.query(  
    Criteria.where("client").is("Chuck Wagon")), Order.class);
```

In this case, the `Criteria` used to construct the `Query` only checks one field. But it can also be used for even more interesting queries. Perhaps you want to get all of Chuck's orders that were placed over the web:

```
List<Order> chucksWebOrders = mongo.find(Query.query(  
    Criteria.where("customer").is("Chuck Wagon")  
        .and("type").is("WEB")), Order.class);
```

And, should you wish to remove a document, the `remove()` method is what you're looking for:

```
mongo.remove(order);
```

As I've said, `MongoOperations` has several methods for working with document data. I encourage you to examine the `JavaDoc` to discover what else you can do with `MongoOperations`.

Typically, you'd inject `MongoOperations` into a repository class of your own design and use its operations to implement the repository methods. But if you don't want to bother writing the repository yourself, then `Spring Data MongoDB` can automatically generate a repository implementation for you at runtime. Let's see how.

12.1.4 Writing a MongoDB repository

To understand how to create repositories with Spring Data MongoDB, let's once again consider what you did in chapter 11 with Spring Data JPA. In listing 11.4, you created an interface named `SpitterRepository` that extends `JpaRepository`. In the same section, you also enabled Spring Data JPA repositories. As a result, Spring Data JPA was able to automatically create an implementation of that interface, including several built-in methods and any methods you added that followed a naming convention.

You've already enabled Spring Data MongoDB repositories with `@EnableMongoRepositories`, so all that's left is to create an interface that you can generate the repository implementation from. Instead of extending `JpaRepository`, however, you'll need to extend `MongoRepository`. The `OrderRepository` interface in the following listing extends `MongoRepository` to provide basic CRUD operations for `Order` documents.

Listing 12.6 Spring Data MongoDB automatically implements repository interfaces

```
package orders.db;
import orders.Order;
import org.springframework.data.mongodb.repository.MongoRepository;

public interface OrderRepository
    extends MongoRepository<Order, String> {
}
```

Because `OrderRepository` extends `MongoRepository`, it transitively extends the `Repository` marker interface. As you'll recall from our exploration of Spring Data JPA, any interface that extends `Repository` will have an implementation automatically generated at runtime. In this case, however, instead of a JPA repository that interacts with a relational database, `OrderRepository` will be implemented to read and write data to a MongoDB document database.

The `MongoRepository` interface has two parameters. The first is the type of `@Document`-annotated object that this repository deals with. The second is the type of the `@Id`-annotated property.

Even though `OrderRepository` doesn't define any methods of its own, it inherits several methods, including several useful methods for CRUD operations on `Order` documents. Table 12.2 describes all of the methods that `OrderRepository` inherits.

The methods in table 12.2 refer to the generic types passed into and returned from the methods. Given that `OrderRepository` extends `MongoRepository<Order, String>`, this means that `T` maps to `Order`, `ID` maps to `String`, and `S` maps to any type that extends `Order`.

Table 12.2 By extending `MongoRepository`, a repository interface inherits several CRUD operations that are automatically implemented by Spring Data MongoDB.

Method	Description
<code>long count();</code>	Returns a count of the documents for the repository type
<code>void delete(Iterable<? extends T>);</code>	Deletes all documents associated with the given objects
<code>void delete(T);</code>	Deletes the document associated with the given object
<code>void delete(ID);</code>	Deletes a document by its ID
<code>void deleteAll();</code>	Deletes all documents for the given repository type
<code>boolean exists(Object);</code>	Returns <code>true</code> if a document associated with the given object exists
<code>boolean exists(ID);</code>	Returns <code>true</code> if a document exists for the given ID
<code>List<T> findAll();</code>	Returns all documents for the repository type
<code>List<T> findAll(Iterable<ID>);</code>	Returns all documents for the given document IDs
<code>List<T> findAll(Pageable);</code>	Returns a paged and sorted list of documents for the repository type
<code>List<T> findAll(Sort);</code>	Returns a sorted list of all documents for the given document ID
<code>T findOne(ID);</code>	Returns a single document for the given ID
<code>save(Iterable<S>);</code>	Saves all documents in the given <code>Iterable</code>
<code>save(S);</code>	Saves a single document for the given object

ADDING CUSTOM QUERY METHODS

The CRUD operations are generally useful, but you may need the repository to provide methods beyond what comes out of the box.

In section 11.3.1, you learned that Spring Data JPA supports a method-naming convention that helps Spring Data to automatically generate implementations for the methods that follow that convention. As it turns out, the very same convention works with Spring Data MongoDB. That means you can add custom methods to `OrderRepository` like this:

```
public interface OrderRepository
    extends MongoRepository<Order, String> {
    List<Order> findByCustomer(String c);
    List<Order> findByCustomerLike(String c);
    List<Order> findByCustomerAndType(String c, String t);
    List<Order> findByCustomerLikeAndType(String c, String t);
}
```


Here you have four new methods, each one finding `Order` objects that match certain criteria. One method finds a list of `Order` where the `customer` property is equal to the value passed into the method. Another finds a list of `Order` where the `customer` property is like the value passed into the method. The next finds `Order` objects whose `customer` and `type` properties are equal to the values passed in. The final method is like the previous, except that the `customer` comparison is a *like* comparison instead of an *equals* comparison.

The `find` query verb is flexible. If you'd prefer, you can use `get` as the query verb:

```
List<Order> getByCustomer(String c);
```

Or if it suits you better, you can use `read`:

```
List<Order> readByCustomer(String c);
```

There's also another special query verb for counting the objects that match:

```
int countByCustomer(String c);
```

As with Spring Data JPA, there's a lot of flexibility in what can go in between the query verb and `By`. For example, you could state what it is you're finding:

```
List<Order> findOrdersByCustomer(String c);
```

There's nothing special about the word `Orders`, though. It has nothing to do with what is being fetched. You could name the method like this:

```
List<Order> findSomeStuffWeNeedByCustomer(String c);
```

You also don't need to return a `List<Order>`. If all you want is a single `Order`, you can simply return `Order`:

```
Order findASingleOrderByCustomer(String c);
```

Here, the first `Order` that would've been found if it were a `List` is what will be returned. If there isn't a match, the method will return `null`.

SPECIFYING QUERIES

As you saw in section 11.3.2, the `@Query` annotation can be used to specify a custom query for a repository method. `@Query` works equally well with MongoDB as it does with JPA. The only material difference is that for MongoDB, `@Query` takes a JSON query string instead of a JPA query.

For example, suppose you need a method that finds all orders of a given type for the customer whose name is "Chuck Wagon". The following method declaration in `OrderRepository` will give you what you need:

```
@Query("{ 'customer': 'Chuck Wagon', 'type' : ?0 }")  
List<Order> findChucksOrders(String t);
```

The JSON given to `@Query` is matched up against all `Order` documents, and any document that matches will be returned. Notice that the `type` property is mapped to `?0`. This indicates that the `type` property should be equal to the zeroth parameter to the

query method. If there were more parameters, they could be referred to with ?1, ?2, and so forth.

MIXING IN CUSTOM REPOSITORY BEHAVIOR

In section 11.3.3, you learned how to mix fully custom methods into an otherwise automatically generated repository. For JPA, that involved creating an intermediary interface that declares the custom method(s), an implementation class for those custom methods, and changing the automatic repository interface to extend the intermediary interface. The steps are the same for a Spring Data MongoDB repository.

Suppose that you need a method that finds all `Order` objects where the document's `type` property matches a given value. You could easily create such a method by giving it a signature of `List<Order> findByType(String t)`. But for the sake of this example, suppose that if the given type is `NET`, then it will query for `Orders` whose `type` is `WEB`. This would be hard to do, even with the `@Query` annotation. A mixin implementation, however, can make it work.

First, define the intermediary interface:

```
package orders.db;
import java.util.List;
import orders.Order;

public interface OrderOperations {
    List<Order> findOrdersByType(String t);
}
```

That's simple enough. Now you can write the mixin implementation. The following listing shows what the implementation might look like.

Listing 12.7 Mixing custom repository functionality into an automatic repository

```
package orders.db;
import java.util.List;
import orders.Order;
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.data.mongodb.core.MongoOperations;
import org.springframework.data.mongodb.core.query.Criteria;
import org.springframework.data.mongodb.core.query.Query;

public class OrderRepositoryImpl implements OrderOperations {
    @Autowired
    private MongoOperations mongo;                                ← Inject MongoOperations

    public List<Order> findOrdersByType(String t) {
        String type = t.equals("NET") ? "WEB" : t;

        Criteria where = Criteria.where("type").is(t);           ← Create query
        Query query = Query.query(where);

        return mongo.find(query, Order.class);                  ← Perform query
    }
}
```

As you can see, the mixin implementation is injected with a `MongoOperations` (the interface that `MongoTemplate` implements). The `findOrdersByType()` method uses the `MongoOperations` to query the database for documents matching the constructed query.

All that's left is to change `OrderRepository` to extend the intermediary `OrderOperations` interface:

```
public interface OrderRepository
    extends MongoRepository<Order, String>, OrderOperations {
    ...
}
```

The thing that ties all of this together is the fact that the implementation class is named `OrderRepositoryImpl`. This is the same name as the `OrderRepository` interface, with an “Impl” suffix. When Spring Data MongoDB generates the repository implementation, it will look for this class and mix it into the automatically generated implementation.

If you don't care for the “Impl” suffix, you can configure Spring Data MongoDB to look for a class with a different suffix in its name. All you need to do is set the `repositoryImplementationPostfix` attribute of `@EnableMongoRepositories` (in the Spring configuration class).

```
@Configuration
@EnableMongoRepositories(basePackages="orders.db",
                        repositoryImplementationPostfix="Stuff")
public class MongoConfig extends AbstractMongoConfiguration {
    ...
}
```

Or, if you're using XML configuration, you can set the `repository-impl-postfix` attribute of `<mongo:repositories>`:

```
<mongo:repositories base-package="orders.db"
                    repository-impl-postfix="Stuff" />
```

Either way, if you configure Spring Data MongoDB this way, it will look for `OrderRepositoryStuff` instead of `OrderRepositoryImpl`.

Document databases such as MongoDB solve a certain class of problems. But just as relational databases aren't a one-size-fits-all solution, neither is MongoDB. And there are certain problems that neither a relational database nor a document database is well-suited for. Fortunately, those aren't the only two choices.

Let's have a look at how Spring Data supports Neo4j, a popular graph database.

12.2 Working with graph data in Neo4j

Whereas document databases store data in coarse-grained documents, graph databases store data in several fine-grained nodes that are connected with each other through relationships. A node in a graph database typically represents a concept in

the database, having properties that describe the state of the node. Relationships connect two nodes and may carry properties of their own.

At their simplest, graph databases are more general purpose than document databases, potentially being a schemaless alternative to relational databases. But because data is structured as a graph, it's possible to traverse relationships to discover things about your data that would be difficult or even impossible with other kinds of databases.

Spring Data Neo4j offers many of the same capabilities as Spring Data JPA and Spring Data MongoDB, albeit targeting the Neo4j graph database. It provides annotations for mapping Java types to nodes and relationships, template-oriented Neo4j access, and automatic generation of repository implementations.

You'll see how to use these features to work with Neo4j. But first, you must configure Spring Data Neo4j.

12.2.1 Configuring Spring Data Neo4j

The key to configuring Spring Data Neo4j is to declare a `GraphDatabaseService` bean and enable automatic Neo4j repository generation. The following listing shows the basic Java configuration needed for Spring Data Neo4j.

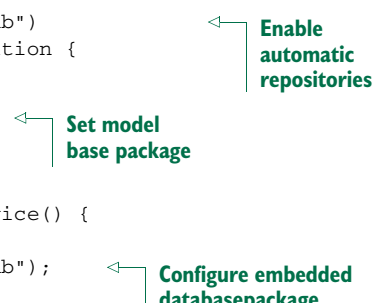
Listing 12.8 Configuring Spring Data Neo4j with `@EnableNeo4jRepositories`

```
package orders.config;
import org.neo4j.graphdb.GraphDatabaseService;
import org.neo4j.graphdb.factory.GraphDatabaseFactory;
import org.springframework.context.annotation.Bean;
import org.springframework.context.annotation.Configuration;
import org.springframework.data.neo4j.config.EnableNeo4jRepositories;
import org.springframework.data.neo4j.config.Neo4jConfiguration;

@Configuration
@EnableNeo4jRepositories(basePackages="orders.db")
public class Neo4jConfig extends Neo4jConfiguration {

    public Neo4jConfig() {
        setBasePackage("orders");
    }

    @Bean(destroyMethod="shutdown")
    public GraphDatabaseService graphDatabaseService() {
        return new GraphDatabaseFactory()
            .newEmbeddedDatabase("/tmp/graphdb");
    }
}
```



Enable automatic repositories

Set model base package

Configure embedded database package

The `@EnableNeo4jRepositories` annotation enables Spring Data Neo4j to automatically generate Neo4j repository implementations. Its `basePackages` is set so that it scans the `orders.db` package for interfaces that extend (directly or indirectly) the marker `Repository` interface.

`Neo4jConfig` extends `Neo4jConfiguration`, which provides convenient methods for configuring Spring Data Neo4j. Among those methods is `setBasePackage()`,

which is called from `Neo4jConfig`'s constructor to tell Spring Data Neo4j that it can find model classes in the `orders` package.

The final piece of the puzzle is to define a `GraphDatabaseService` bean. In this case, the `graphDatabaseService()` method uses `GraphDatabaseFactory` to create an embedded Neo4j database. With Neo4j, an embedded database shouldn't be confused with an in-memory database. "Embedded" means that the database engine is running within the same JVM as a part of your application rather than as a separate server. The data is still persisted to the filesystem (at `/tmp/graphdb` in this case).

Alternatively, you might want to configure a `GraphDatabaseService` that references a remote Neo4j server. If you have the `spring-data-neo4j-rest` library in your application's classpath, you can configure `SpringRestGraphDatabase`, which accesses a remote Neo4j database over a RESTful API:

```
@Bean(destroyMethod="shutdown")
public GraphDatabaseService graphDatabaseService() {
    return new SpringRestGraphDatabase(
        "http://graphdbserver:7474/db/data/");
}
```

As shown here, `SpringRestGraphDatabase` is configured to assume that the remote database doesn't require authentication. In a production setting, however, it's likely that you'll want to secure the database server. In that case, you'll want to provide your application's credentials when creating the `SpringRestGraphDatabase`:

```
@Bean(destroyMethod="shutdown")
public GraphDatabaseService graphDatabaseService(Environment env) {
    return new SpringRestGraphDatabase(
        "http://graphdbserver:7474/db/data/",
        env.getProperty("db.username"), env.getProperty("db.password"));
}
```

Here, the credentials are obtained via the injected `Environment` to avoid hard-coding them in the configuration class.

Spring Data Neo4j also offers an XML configuration namespace. If you'd rather configure Spring Data Neo4j in XML, you can use the `<neo4j:config>` and `<neo4j:repositories>` elements from that namespace. Listing 12.9 shows an XML configuration that's equivalent to the Java configuration in listing 12.8.

Listing 12.9 Spring Data Neo4j can also be configured in XML.

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
    xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
    xmlns:neo4j="http://www.springframework.org/schema/data/neo4j"
    xsi:schemaLocation="
        http://www.springframework.org/schema/beans
        http://www.springframework.org/schema/beans/spring-beans.xsd
        http://www.springframework.org/schema/data/neo4j
        http://www.springframework.org/schema/data/neo4j/spring-neo4j.xsd">
```

```

<neo4j:config
  storeDirectory="/tmp/graphdb"
  base-package="orders" />
<neo4j:repositories base-package="orders.db" />
</beans>

```

← Configure Neo4j database details

← Enable repository generation

The `<neo4j:config>` element configures the details of how to access the database. In this case, it configures Spring Data Neo4j to work with an embedded database. Specifically, the `storeDirectory` attribute specifies the path in the filesystem where the data will be persisted. The `base-package` attribute sets the package where the model classes are defined.

As for `<neo4j:repositories>`, it enables Spring Data Neo4j to automatically generate repository implementations by scanning the `orders.db` package and looking for interfaces that extend the `Repository` interface.

To configure Spring Neo4j to access a remote Neo4j server, all you need to do is declare a `SpringRestGraphDatabase` bean and set `<neo4j:config>`'s `graphDatabaseService` attribute:

```

<neo4j:config base-package="orders"
  graphDatabaseService="graphDatabaseService" />
<bean id="graphDatabaseService" class=
  "org.springframework.data.neo4j.rest.SpringRestGraphDatabase">
  <constructor-arg value="http://graphdbserver:7474/db/data/" />
  <constructor-arg value="db.username" />
  <constructor-arg value="db.password" />
</bean>

```

Whether you configure Spring Data Neo4j using Java or XML, you'll need to be sure that the domain types are under a package specified as a base package (`@EnableNeo4jRepositories`'s `basePackages` attribute or `<neo4j:config>`'s `base-package` attribute). They'll also need to be annotated as either node entities or relationship entities. That's what you'll do next.

12.2.2 Annotating graph entities

Neo4j defines two kinds of entities: nodes and relationships. Node entities typically represent the things in your application, whereas relationship entities define how those things are related.

Spring Data Neo4j provides several annotations that you can apply to domain types and their fields for persistence in Neo4j. Table 12.3 describes these annotations.


To see how a few of these annotations are used, let's apply these annotations to our `order/item` example.

One way you can model the data is to designate an order as a node that's related to one or more items. Figure 12.2 illustrates this model as a graph.

To designate orders as nodes, you'll need to annotate the `Order` class with `@NodeEntity`. The next listing shows the `Order` class annotated with `@NodeEntity`, as well as a few other annotations from table 12.3.

Table 12.3 Spring Data Neo4j annotations you can use to map domain types to nodes and relationships in a graph

Annotation	Description
@NodeEntity	Declares a Java type as a node entity
@RelationshipEntity	Declares a Java type as a relationship entity
@StartNode	Declares a property as the start node of a relationship entity
@EndNode	Declares a property as the end node of a relationship entity
@Fetch	Declares a property on an entity to be eagerly loaded
@GraphId	Declares a property as the ID field of an entity (the field must be a Long)
@GraphProperty	Explicitly declares a property
@GraphTraversal	Declares a property to automatically provide an iterable that's built by following a graph traversal
@Indexed	Declares a property to be indexed
@Labels	Declares the labels for an @NodeEntity
@Query	Declares a property to automatically provide an iterable that's built by executing a given Cypher query
@QueryResult	Declares a Java class or interface as being able to hold the results of a query
@RelatedTo	Declares a simple relationship between the current @NodeEntity and another @NodeEntity via a property
@RelatedToVia	Declares a field on an @NodeEntity as referencing an @RelationshipEntity that the node belongs to
@RelationshipType	Declares a field as the type of a relationship entity
@ResultColumn	Declares a property on an @QueryResult-annotated type to capture a specific field from a query result

**Figure 12.2** A simple relationship connects two nodes but carries no properties of its own.**Listing 12.10** Order is annotated to be a node in the graph database

```

package orders;
import java.util.LinkedHashSet;
import java.util.Set;
import org.springframework.data.neo4j.annotation.GraphId;
  
```

```

import org.springframework.data.neo4j.annotation.NodeEntity;
import org.springframework.data.neo4j.annotation.RelatedTo;

@NodeEntity                                     ← Orders are nodes
public class Order {

    @GraphId                                     ← The graph ID
    private Long id;
    private String customer;
    private String type;

    @RelatedTo(type="HAS_ITEMS")                ← Relationship to items
    private Set<Item> items = new LinkedHashSet<Item>();

    ...
}

```

In addition to the `@NodeEntity` at the class level, notice that the `id` property is annotated with `@GraphId`. All entities in Neo4j must have a graph ID. This is roughly analogous to the `@Id`-annotated properties of a JPA `@Entity` or a MongoDB `@Document`. It's required that the `@GraphId`-annotated property be a `Long`.

The `customer` and `type` properties remain annotation-free. Unless they are transient, they'll be properties on the node in the database.

The `items` property is annotated with `@RelatedTo`, indicating that an `Order` is related to a `Set` of `Item`. The `type` attribute essentially labels the relationship. It can be given any value, but it's commonly given human-readable text that briefly describes the nature of the relationship. Later you'll use this label in queries to query across relationships.

As for the `Item` class itself, the following listing shows how it's annotated for graph persistence.

Listing 12.11 Items are also represented as nodes in the graph database.

```

package orders;
import org.springframework.data.neo4j.annotation.GraphId;
import org.springframework.data.neo4j.annotation.NodeEntity;

@NodeEntity                                     ← Items are nodes
public class Item {

    @GraphId                                     ← The graph ID
    private Long id;
    private String product;
    private double price;
    private int quantity;

    ...
}

```

As with `Order`, `Item` is annotated as `@NodeEntity` to designate it as a node. It also has a `Long` property annotated to be the node's graph ID with `@GraphId`. The `product`,

price, and quantity properties will also be persisted as node properties in the graph database.

The relationship between `Order` and `Item` is simple in that it doesn't carry any data of its own. Therefore, the `@RelatedTo` annotation is sufficient to define the relationship. But not all relationships are so simple.

Let's reconsider how we've modeled this data to see how to work with more complex relationships. In the current data model, we've combined the concepts of a line item and a product into the `Item` class. When you think about it, however, an order is related to one or more products. The relationship between an order and a product constitutes a line item of the order. Figure 12.3 illustrates an alternative way to model the data in a graph.




Figure 12.3 A relationship entity is a relationship that has properties of its own.

In this new model, the quantity of products in the order is a property of the line item, and a product is a different concept. As before, orders are nodes and so are products. Line items are relationships. But now that a line item must carry a quantity value, the relationship can't be simple. You're going to need to define a class that represents a line item, such as `LineItem` in the next listing.

Listing 12.12 A `LineItem` connects an `Order` node and a `Product` node.

```

package orders;
import org.springframework.data.neo4j.annotation.EndNode;
import org.springframework.data.neo4j.annotation.GraphId;
import org.springframework.data.neo4j.annotation.RelationshipEntity;
import org.springframework.data.neo4j.annotation.StartNode;

@RelationshipEntity(type="HAS_LINE_ITEM_FOR")  <— LinelItem is a relationship
public class LineItem {

    @GraphId                                  <— The graph ID
    private Long id;

    @StartNode                                <— The start node
    private Order order;

    @EndNode                                  <— The end node
    private Product product;

    private int quantity;

    ...
}
  
```

Whereas `Order` was annotated with `@NodeEntity` to designate it as a node, `LineItem` is annotated with `@RelationshipEntity`. `LineItem` also has an `id` property annotated

with `@GraphId`. Again, all entities, both node entities and relationship entities, must have a graph ID and it must be of type `Long`.

What makes relationship entities special is that they connect two node entities. The `@StartNode` and `@EndNode` annotations are applied to properties that define each end of a relationship. In this case, the `Order` is the start node and the `Product` is the end node.

Finally, `LineItem` has a `quantity` property that will be persisted to the database when the relationship is created.

Now that the domain is annotated, you're ready to start saving and reading nodes and relationships. We'll start by looking at how you can use Spring Data Neo4j's template-oriented data access with `Neo4jTemplate`.

12.2.3 Working with Neo4jTemplate

Just as Spring Data MongoDB provides `MongoTemplate` for template-based MongoDB persistence, Spring Data Neo4j brings `Neo4jTemplate` to the table to work with nodes and relationships in the Neo4j graph database. If you've configured Spring Data Neo4j as shown earlier, then there's already a `Neo4jTemplate` bean in the Spring application context. All you need to do is inject it wherever you need it.

For example, you might autowire it directly into a bean property:

```
@Autowired
private Neo4jOperations neo4j;
```

`Neo4jTemplate` defines several dozen methods, including methods for saving nodes, deleting nodes, and creating relationships between nodes. There's not enough space to cover all of them, but let's have a look at a few of the most commonly used methods that `Neo4jTemplate` provides.

One of the first and most basic things you might want to do with `Neo4jTemplate` is to save an object as a node. Assuming that the object is annotated with `@NodeEntity`, you can use the `save()` method like this:

```
Order order = ...;
Order savedOrder = neo4j.save(order);
```

If you happen to know the object's graph ID, you can fetch it using the `findOne()` method:

```
Order order = neo4j.findOne(42, Order.class);
```

If there is no node with the given ID, then `findOne()` will throw a `NotFoundException`.

If you'd like to retrieve all objects of a given type, you can use the `findAll()` method:

```
EndResult<Order> allOrders = neo4j.findAll(Order.class);
```

The `EndResult` returned here is an `Iterable`, enabling it to be used in for-each looping and anywhere else an `Iterable` may be used. If no such nodes exist, `findAll()` will return an empty `Iterable`.

If all you need to know is a count of how many objects of a given type are in the Neo4j database, you can call the `count()` method:

```
long orderCount = count(Order.class);
```

The `delete()` method can be used to delete an object:

```
neo4j.delete(order);
```

One of the most interesting methods provided by `Neo4jTemplate` is the `createRelationshipBetween()` method. As you might guess, it creates a relationship between two nodes. For example, you could create a `LineItem` relationship between an `Order` node and a `Product` node:

```
Order order = ...;
Product prod = ...;
LineItem lineItem = neo4j.createRelationshipBetween(
    order, prod, LineItem.class, "HAS_LINE_ITEM_FOR", false);
lineItem.setQuantity(5);
neo4j.save(lineItem);
```

The first two parameters to `createRelationshipBetween()` are the objects whose nodes will be at each end of the relationship. The next parameter specifies the `@RelationshipEntity`-annotated type that will represent the relationship. Next, you specify a `String` value that describes the nature of the relationship. The final parameter is a boolean that indicates whether or not duplicate relationships are allowed between the two node entities.

`createRelationshipBetween()` returns an instance of the relationship class. From there you can set any properties you'd like. The preceding example sets the `quantity` property. When you're done, you call `save()` to save the relationship to the database.

The `Neo4jTemplate` offers a straightforward way to work with nodes and relationships in a Neo4j graph database. But it requires that you write your own repository implementations that delegate to `Neo4jTemplate`. Let's see how Spring Data Neo4j can automatically generate repository implementations for you.

12.2.4 Creating automatic Neo4j repositories

One of the most awesome things that most Spring Data projects do is automatically generate implementations for a repository interface. You've already seen this with Spring Data JPA and Spring Data MongoDB. Not to be left out, Spring Data Neo4j also supports automatic repository generation.

You've already added `@EnableNeo4jRepositories` to your configuration, so Spring Data Neo4j is already set to generate repositories. All you need to do is write the interfaces. The following `OrderRepository` interface is a good start:

```
package orders.db;
import orders.Order;
import org.springframework.data.neo4j.repository.GraphRepository;

public interface OrderRepository extends GraphRepository<Order> {}
```

Just like the other Spring Data projects, Spring Data Neo4j triggers repository generation for interfaces that extend the `Repository` interface. In this case, `OrderRepository` extends `GraphRepository`, which indirectly extends `Repository`. Therefore, Spring Data Neo4j will generate an implementation for `OrderRepository` at runtime.

Notice that `GraphRepository` is parameterized with `Order`, the type of entity that the repository works with. Because Neo4j requires that graph IDs be of type `Long`, there's no need to specify the ID type when extending `GraphRepository`.

Out of the box, you get several common CRUD operations, much like what `JpaRepository` and `MongoRepository` provide. Table 12.4 describes the methods you get by extending `GraphRepository`.

Table 12.4 By extending `GraphRepository`, a repository interface inherits several CRUD operations that are automatically implemented by Spring Data Neo4j.

Method	Description
<code>long count();</code>	Returns a count of how many entities of the target type are in the database
<code>void delete(Iterable<? extends T>);</code>	Deletes several entities
<code>void delete(Long id);</code>	Deletes a single entity given its ID
<code>void delete(T);</code>	Deletes a single entity
<code>void deleteAll();</code>	Deletes all entities of the target type
<code>boolean exists(Long id);</code>	Checks for the existence of an entity given its ID
<code>EndResult<T> findAll();</code>	Retrieves all entities of the target type
<code>Iterable<T> findAll(Iterable<Long>);</code>	Retrieves all entities of the target type for the given IDs
<code>Page<T> findAll(Pageable);</code>	Retrieves a paged and sorted list of all entities of the target type
<code>EndResult<T> findAll(Sort);</code>	Retrieves a sorted list of all entities of the target type
<code>EndResult<T> findAllBySchemaPropertyValue(String, Object);</code>	Retrieves all entities where a given property matches the given value
<code>Iterable<T> findAllByTraversal(N, TraversalDescription);</code>	Retrieves all entities obtained by following a graph traversal starting at a given node
<code>T findBySchemaPropertyValue(String, Object);</code>	Finds a single entity where a given property matches a given value
<code>T findOne(Long);</code>	Finds a single entity given its ID
<code>EndResult<T> query(String, Map<String, Object>);</code>	Finds all entities that match a given Cypher query

Table 12.4 By extending `GraphRepository`, a repository interface inherits several CRUD operations that are automatically implemented by Spring Data Neo4j. (continued)

Method	Description
<code>Iterable<T> save(Iterable<T>);</code>	Saves several entities
<code>S save(S);</code>	Saves a single entity

There's not enough space to cover all of these methods, but there are a few methods that you'll get a lot of use out of. For example, the following line saves a single `Order` entity:

```
Order savedOrder = orderRepository.save(order);
```

When the entity is saved, the `save()` method returns the saved entity, which now should have its `@GraphId`-annotated property populated if it was `null` before.

You can look up a single entity by calling the `findOne()` method. For example, this line will look up an `Order` whose graph ID is 4:

```
Order order = orderRepository.findOne(4L);
```

Or, you can look up all orders:

```
EndResult<Order> allOrders = orderRepository.findAll();
```

Of course, you may want to delete an entity. In that case, you can use the `delete()` method:

```
delete(order);
```

This will delete the given `Order` node from the database. If you only have the graph ID, you can pass it to `delete()` instead of the node type itself:

```
delete(orderId);
```

If you want to do custom queries, you could use the `query()` method to execute an arbitrary Cypher query against the graph. But that's not much different than working with the `query()` method from `Neo4jTemplate`. Instead, you can add your own query methods to `OrderRepository`.

ADDING QUERY METHODS

You've already seen how to add query methods that follow a naming convention using Spring Data JPA and Spring Data MongoDB. It'd be awfully disappointing if Spring Data Neo4j didn't offer the same capability.

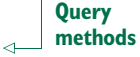
As the next listing shows, there's no need to be disappointed.

Listing 12.13 Defining query methods by following a naming convention.

```
package orders.db;
import java.util.List;
import orders.Order;
```

```
import org.springframework.data.neo4j.repository.GraphRepository;

public interface OrderRepository extends GraphRepository<Order> {
    List<Order> findByCustomer(String customer);
    List<Order> findByCustomerAndType(String customer, String type);
}
```



Here you add two query methods. One finds all `Order` nodes where the `customer` property is equal to the given `String` value. The other method is similar, but in addition to matching the `customer` property, the `Order` nodes must also have a `type` property equal to the given `type`.

We've already discussed the naming convention for query methods, so there's no need to dwell on it any further. Refer to the previous chapter's discussion of Spring Data JPA for a refresher on how to write these methods.

SPECIFYING CUSTOM QUERIES

When the naming convention doesn't meet your needs, you also have the option of annotating a method with `@Query` to specify your own query. You've seen `@Query` before. With Spring Data JPA, you used it to specify a JPA query for a repository method. With Spring Data MongoDB, you used it to specify a JSON-matching query. When using Spring Data Neo4j, however, you must specify a Cypher query:

```
@Query("match (o:Order)-[:HAS_ITEMS]->(i:Item) " +
        "where i.product='Spring in Action' return o")
List<Order> findSiAOrders();
```

Here, the `findSiAOrders()` is annotated with `@Query` and given a Cypher query to find all `Order` nodes that are related to an `Item` whose `product` property is equal to "Spring in Action".

MIXING IN CUSTOM REPOSITORY BEHAVIOR

When neither the naming convention nor the `@Query` methods meet your needs, you always have the option of mixing in custom repository logic.

For example, suppose that you want to write the implementation of `findSiAOrders()` yourself, instead of relying on the `@Query` annotation. You can start by defining an intermediary interface that carries the definition of the `findSiAOrders()` method:

```
package orders.db;
import java.util.List;
import orders.Order;

public interface OrderOperations {
    List<Order> findSiAOrders();
}
```

Then you can change `OrderRepository` to extend `OrderOperations` in addition to `GraphRepository`:

```
public interface OrderRepository
    extends GraphRepository<Order>, OrderOperations {

    ...

}
```

Finally, you need to write the implementation itself. As with Spring Data JPA and Spring Data MongoDB, Spring Data Neo4j will look for an implementation class whose name is the same as the repository interface with an “Impl” suffix. Therefore, you need to create an `OrderRepositoryImpl` class. The following listing shows `OrderRepositoryImpl`, which implements `findSiAOrders()`.

Listing 12.14 Mixing custom functionality into `OrderRepository`

```
package orders.db;
import java.util.Collections;
import java.util.List;
import java.util.Map;
import orders.Order;
import org.neo4j.helpers.collection.IteratorUtil;
import org.springframework.beans.factory.annotation.Autowired;
import org.springframework.data.neo4j.conversion.EndResult;
import org.springframework.data.neo4j.conversion.Result;
import org.springframework.data.neo4j.template.Neo4jOperations;

public class OrderRepositoryImpl implements OrderOperations {
    private final Neo4jOperations neo4j;

    @Autowired
    public OrderRepositoryImpl(Neo4jOperations neo4j) {
        this.neo4j = neo4j;
    }

    public List<Order> findSiAOrders() {
        Result<Map<String, Object>> result = neo4j.query(
            "match (o:Order)-[:HAS_ITEMS]->(i:Item) " +
            "where i.product='Spring in Action' return o",
            EndResult<Order> endResult = result.to(Order.class);
        return IteratorUtil.asList(endResult);
    }
}
```

Implement intermediate interface

Inject Neo4jOperations

Perform query

Convert to EndResult<Order>

Convert to a List<Order>

`OrderRepositoryImpl` is injected with `Neo4jOperations` (specifically, an instance of `Neo4jTemplate`), which it uses to query the database. Because the `query()` method returns a `Result<Map<String, Object>>`, you’ll need to convert it to a `List<Order>`. The first step is to call the `to()` method on the `Result` to produce an `EndResult<Order>`. Then you use Neo4j’s `IteratorUtil.asList()` to convert the `EndResult<Order>` to a `List<Order>`, which is then returned.

Graph databases such as Neo4j are wonderful for capturing data that’s represented well as nodes and relationships. When you consider that the world we live in is made

up of all kinds of things that are related to each other, graph databases may be suitable for a wide range of domains. And, speaking personally, I'll admit that I have a strong fondness for Neo4j.

But sometimes your data needs are simpler. Sometimes you only need to store a value somewhere and be able to retrieve it later with a key. Let's see how Spring Data enables key-value data persistence using the Redis key-value store.

12.3 Working with key-value data in Redis

Redis is a special kind of database known as a key-value store. As the name implies, key-value stores keep key-value pairs. In fact, key-value stores share a lot in common with hash maps. To call them persistent hash maps would not be too great of an oversimplification.

When you think about it, there aren't too many kinds of queries that you can perform against a hash map ... or a key-value store. You can store a value at a particular key, and you can fetch the value for a particular key. That's about it. Consequently, Spring Data's automatic repository support doesn't make a lot of sense when applied to Redis. On the other hand, Spring Data's other key feature, template-oriented data access, can come in handy when working with Redis.

Spring Data Redis comes with a couple of template implementations for storing data to and fetching it from a Redis database. You'll see how to use them soon. But to create one of Spring Data Redis's templates, you'll need a Redis connection factory. Fortunately, Spring Data Redis offers four to choose from.

12.3.1 Connecting to Redis

A Redis connection factory produces connections to a Redis database server. Spring Data Redis comes with connection factories for four Redis client implementations:

- JedisConnectionFactory
- JredisConnectionFactory
- LettuceConnectionFactory
- SrpConnectionFactory

The choice is up to you. I encourage you to do your own testing and benchmarking to determine which Redis client and connection factory fits your needs best. From Spring Data Redis's perspective, all of these connection factories are equally suitable.

Once you've made your choice, you can configure the connection factory as a bean in Spring. For example, here's how you might configure the JedisConnectionFactory bean:

```
@Bean
public RedisConnectionFactory redisCF() {
    return new JedisConnectionFactory();
}
```


Instantiating the connection factory via its default constructor results in a connection factory that creates its connections for localhost, port 6379, and with no password. If your Redis server is running on a different host or port, you can set those properties when you create the connection factory:

```
@Bean
public RedisConnectionFactory redisCF() {
    JedisConnectionFactory cf = new JedisConnectionFactory();
    cf.setHostName("redis-server");
    cf.setPort(7379);
    return cf;
}
```

Similarly, if your Redis server is configured to require authorization from clients, you can set the password by calling `setPassword()`:

```
@Bean
public RedisConnectionFactory redisCF() {
    JedisConnectionFactory cf = new JedisConnectionFactory();
    cf.setHostName("redis-server");
    cf.setPort(7379);
    cf.setPassword("foobared");
    return cf;
}
```

In all of these examples, I've been assuming the `JedisConnectionFactory`. If you've made a different choice, then any of the other connection factories can be a drop-in replacement. For example, if you'd rather use the `LettuceConnectionFactory`, you can configure it like this:

```
@Bean
public RedisConnectionFactory redisCF() {
    JedisConnectionFactory cf = new LettuceConnectionFactory();
    cf.setHostName("redis-server");
    cf.setPort(7379);
    cf.setPassword("foobared");
    return cf;
}
```

All of the Redis connection factories have `setHostName()`, `setPort()`, and `setPassword()` methods. This makes them virtually identical in terms of configuration.

Now that you have a Redis connection factory, you're ready to start working with Spring Data Redis's templates.

12.3.2 Working with RedisTemplate

As their names suggest, the Redis connection factories produce connections (as `RedisConnection`) to a Redis key-value store. Using `RedisConnection`, you can store and read data. For example, you might obtain a connection and use it to store a greeting like this:

```
RedisConnectionFactory cf = ...;
RedisConnection conn = cf.getConnection();
```

```
conn.set("greeting".getBytes(), "Hello World".getBytes());
```

Likewise, you could retrieve that greeting value using a `RedisConnection` like this:

```
byte[] greetingBytes = conn.get("greeting".getBytes());
String greeting = new String(greetingBytes);
```

No doubt, this will work. But do you really like working with arrays of bytes?

As with other Spring Data projects, Spring Data Redis offers a higher-level data access option with templates. In fact, Spring Data Redis offers two templates:

- `RedisTemplate`
- `StringRedisTemplate`

`RedisTemplate` is a class that greatly simplifies Redis data access, enabling you to persist keys and values of any type, not just byte arrays. In recognition of the fact that keys and values are frequently Strings, `StringRedisTemplate` extends `RedisTemplate` to have a String focus.

Assuming you have a `RedisConnectionFactory` available, you can construct a `RedisTemplate` like this:

```
RedisConnectionFactory cf = ...;
RedisTemplate<String, Product> redis =
    new RedisTemplate<String, Product>();
redis.setConnectionFactory(cf);
```

Note that `RedisTemplate` is parameterized with two types. The first type is that of the key, and the second is that of the value. In the `RedisTemplate` constructed here, `Product` objects will be stored as values assigned to String keys.

If you know that you'll be working with both String values and String keys, then you should consider using `StringRedisTemplate` instead of `RedisTemplate`:

```
RedisConnectionFactory cf = ...;
StringRedisTemplate redis = new StringRedisTemplate(cf);
```

Notice that, unlike `RedisTemplate`, `StringRedisTemplate` has a constructor that accepts a `RedisConnectionFactory`. Therefore, there's no need to call `setConnectionFactory()` after construction.

Although it's not required, if you'll be using `RedisTemplate` or `StringRedisTemplate` frequently, you might consider configuring them as beans to be injected where they're needed. Here's a simple `@Bean` method declaring a `RedisTemplate` bean:

```
@Bean
public RedisTemplate<String, Product>
    redisTemplate(RedisConnectionFactory cf) {
    RedisTemplate<String, Product> redis =
        new RedisTemplate<String, Product>();
    redis.setConnectionFactory(cf);
    return redis;
}
```

And here's a bean method to declare a `StringRedisTemplate` bean:

```
@Bean
```

```
public StringRedisTemplate
    stringRedisTemplate(RedisConnectionFactory cf) {
    return new StringRedisTemplate(cf);
}
```

Once you have a `RedisTemplate` (or `StringRedisTemplate`), you can start saving, fetching, and deleting key-value entries. Most of the operations provided by `RedisTemplate` are available via the sub-APIs listed in table 12.5.

Table 12.5 `RedisTemplate` offers much of its functionality via sub-APIs, which differentiate single values from collection values.

Method	Sub-API interface	Description
<code>opsForValue()</code>	<code>ValueOperations<K, V></code>	Operations for working with entries having simple values
<code>opsForList()</code>	<code>ListOperations<K, V></code>	Operations for working with entries having list values
<code>opsForSet()</code>	<code>SetOperations<K, V></code>	Operations for working with entries having set values
<code>opsForZSet()</code>	<code>ZSetOperations<K, V></code>	Operations for working with entries having ZSet (sorted set) values
<code>opsForHash()</code>	<code>HashOperations<K, HK, HV></code>	Operations for working with entries having hash values
<code>boundValueOps(K)</code>	<code>BoundValueOperations<K,V></code>	Operations for working with simple values bound to a given key
<code>boundListOps(K)</code>	<code>BoundListOperations<K,V></code>	Operations for working with list values bound to a given key
<code>boundSetOps(K)</code>	<code>BoundSetOperations<K,V></code>	Operations for working with set values bound to a given key
<code>boundZSet(K)</code>	<code>BoundZSetOperations<K,V></code>	Operations for working with ZSet (sorted set) values bound to a given key
<code>boundHashOps(K)</code>	<code>BoundHashOperations<K,V></code>	Operations for working with hash values bound to a given key

As you can see, the sub-APIs in table 12.5 are available through methods on `RedisTemplate` (and `StringRedisTemplate`). Each one provides operations that work with entries based on whether the value is a simple value or a collection of values.

Across all of these sub-APIs, there are several dozen methods for saving and fetching data in Redis. We don't have space enough to cover them all, but we'll look at a handful of the most common operations you'll need.

WORKING WITH SIMPLE VALUES

Suppose that you want to save a `Product` to a `RedisTemplate<String, Product>` where the key is the value of the `sku` property. The following snippet of code will do that via `opsForValue()`:

```
redis.opsForValue().set(product.getSku(), product);
```

Similarly, if you wanted to fetch a product whose `sku` is `123456`, you could use this snippet:

```
Product product = redis.opsForValue().get("123456");
```

If no entry can be found with the given key, `null` will be returned.

WORKING WITH LISTS

Working with list values is similarly straightforward via `opsForList()`. For example, you can add a value to the end of a list entry like this:

```
redis.opsForList().rightPush("cart", product);
```

This adds a `Product` to the end of the list stored at the key `cart`. If a list doesn't already exist at that key, one will be created.

Whereas the `rightPush()` method adds an element to the end of a list entry, `leftPush()` inserts a value at the beginning:

```
redis.opsForList().leftPush("cart", product);
```

There are a number of ways you can fetch an item from a list. You can pop an entry off of either end using `leftPop()` or `rightPop()`:

```
Product first = redis.opsForList().leftPop("cart");  
Product last = redis.opsForList().rightPop("cart");
```

Aside from fetching a value from the list, these two pop methods have the side effect of removing the popped items from the list. If you'd rather simply retrieve the value (perhaps even from the middle of the list), you can use the `range()` method:

```
List<Product> products = redis.opsForList().range("cart", 2, 12);
```

The `range()` method doesn't remove any values from the list entry, but it does retrieve one or more values given the key and a range of indexes. The preceding example fetches as many as eleven entries starting with the entry at index 2 and going through index 12 (inclusive). If the range exceeds the bounds of the list, then only the entries within those indexes will be returned. If no entries fall within the indexes, an empty list will be returned.

PERFORMING OPERATIONS ON SETS

In addition to lists, you can also work with sets via the `opsForSet()` method. The most basic thing you can do is add an item to a set entry:

```
redis.opsForSet().add("cart", product);
```

Once you have a few set entries created and populated with values, you can perform interesting operations against those sets, including difference, intersection, and union.

```
List<Product> diff = redis.opsForSet().difference("cart1", "cart2");
List<Product> union = redis.opsForSet().union("cart1", "cart2");
List<Product> isect = redis.opsForSet().isect("cart1", "cart2");
```

Of course, you can also remove items:

```
redis.opsForSet().remove(product);
```

And you can even fetch a random element from the set:

```
Product random = redis.opsForSet().randomMember("cart");
```

As sets don't have indexes or any implicit ordering, you can't pinpoint and fetch a single item from the set.

BINDING TO A KEY

Table 12.5 includes five sub-APIs for working with operations bound to a given key. These sub-APIs mirror the other sub-APIs, but focus on a given key.

As an example of how these are used, let's consider the case where you're storing `Product` objects in a list entry whose key is `cart`. In that scenario, suppose that you want to pop an item from the right end of the list and then add three new items to the end of the list. You can do that using the `BoundListOperations` returned from calling `boundListOps()`:

```
BoundListOperations<String, Product> cart =
    redis.boundListOps("cart");
Product popped = cart.rightPop();
cart.rightPush(product1);
cart.rightPush(product2);
cart.rightPush(product3);
```

Notice that the only time that the entry's key is mentioned is when calling `boundListOps()`. All of the operations performed against the returned `BoundListOperations` will be applied to that entry.

12.3.3 Setting key and value serializers

When an entry is saved to the Redis key-value store, both the key and the value are serialized using a Redis serializer. Spring Data Redis comes with several such serializers, including these:

- `GenericToStringSerializer`—Serializes using a Spring conversion service
- `JacksonJsonRedisSerializer`—Serializes objects to JSON using Jackson 1
- `Jackson2JsonRedisSerializer`—Serializes objects to JSON using Jackson 2
- `JdkSerializationRedisSerializer`—Uses Java serialization
- `OxmSerializer`—Serializes using marshalers and unmarshalers from Spring's O/X mapping, for XML serialization

- `StringRedisSerializer`—Serializes `String` keys and values

All of these serializers implement the `RedisSerializer` interface, so if there's not one to suit your needs, you can always create your own serializer.

`RedisTemplate` uses `JdkSerializationRedisSerializer`, which means that keys and values are serialized through Java. As you might expect, `StringRedisTemplate` uses `StringRedisSerializer` by default, essentially converting the `String` values to and from byte arrays. These defaults are suitable for many cases, but you may find it helpful to plug in a different serializer.

For example, suppose that when using `RedisTemplate`, you want to serialize `Product` values to JSON with `String` keys. The `setKeySerializer()` and `setValueSerializer()` methods of `RedisTemplate` are what you need:

```
@Bean
public RedisTemplate<String, Product>
    redisTemplate(RedisConnectionFactory cf) {
    RedisTemplate<String, Product> redis =
        new RedisTemplate<String, Product>();
    redis.setConnectionFactory(cf);
    redis.setKeySerializer(new StringRedisSerializer());
    redis.setValueSerializer(
        new Jackson2JsonRedisSerializer<Product>(Product.class));
    return redis;
}
```

Here, you set `RedisTemplate` to always use `StringRedisSerializer` when serializing key values. You also specify that it should use `Jackson2JsonRedisSerializer` only when serializing `Product` values.

12.4 Summary

Gone are the days when the only choice for data persistence was a relational database. Now there are several different kinds of databases, each representing data in different forms and offering capabilities to suit a variety of domain models. The Spring Data project enables developers to use these databases in their Spring applications and to use abstractions that are reasonably consistent across the various database choices.

In this chapter, we built on what you learned about Spring Data in the previous chapter when using JPA, applying it to the MongoDB document database and the Neo4j graph database. Just like their JPA counterpart, the Spring Data MongoDB and Spring Data Neo4j projects both offer automatic generation of repositories based on interface definitions. Additionally, you saw how to use the annotations provided by the Spring Data projects to map domain types to documents, nodes, and relationships.

Spring Data also enables data to be persisted to the Redis key-value store. Key-value stores are significantly simpler and thus do not require support for automatic repositories or mapping annotations. Nevertheless, Spring Data Redis offers two different template classes for working with the Redis key-value store.

No matter what kind of database you choose, fetching data from the database is a costly operation. In fact, database queries are often the biggest performance bottlenecks in any application. Now that you've seen how to store and fetch data from a variety of data sources, let's look at how to avoid that bottleneck. In the next chapter, you'll see how to apply declarative caching to prevent unnecessary database fetches.

Spring IN ACTION FOURTH EDITION

Craig Walls



Designed in 2003 as a lighter approach to J2EE development, Spring Framework has since become a standard choice for building enterprise applications and required knowledge for Java developers. Spring 4, the latest major version, provides full Java 8 integration along with key upgrades like new annotations for the IoC container, improvements to Spring Expression Language, and much-needed support for REST. Whether you're just discovering Spring or you want to absorb the new features, there's no better way to master Spring than with this book.

Spring in Action, Fourth Edition is a hands-on guide to the Spring Framework. It covers Spring core, along with the latest updates to Spring MVC, Security, Web Flow, and more. You'll move between short snippets and an ongoing example as you learn to build simple and efficient JEE applications. Author Craig Walls has a special knack for crisp and entertaining examples that zoom in on the features and techniques you really need.

What's Inside

- Updated for Spring 4
- Spring Data for NoSQL
- Simplify configuration with annotations and definition profiles
- Working with RESTful resources

Nearly 100,000 developers have used this book to learn Spring! It requires a working knowledge of Java.

Craig Walls is a software developer at Pivotal. He's a popular author and a frequent speaker at user groups and conferences. Craig lives in Cross Roads, Texas.

To download their free eBook in PDF, ePub, and Kindle formats, owners of this book should visit manning.com/SpringinActionFourthEdition

“The best book for Spring—updated and revised.”

—Gregor Zurowski, Sotheby's

“The classic, remastered and full of awesomeness.”

—Mario Arias, Cake Solutions Ltd

“Informative, accurate and insightful!”

—Jeelani Shaik, D3Banking.com

“After ten years, this is still the clearest and most comprehensive introduction to the core concepts of the Spring platform.”

—James Wright, Sword-Apak

